

A computer-vision based training coach for computerized physical training

George Davies^[0009–0008–4132–5676]

University of Lincoln, UK 27421138@students.lincoln.ac.uk

1 Introduction

In many stories of science-fiction, there are robots that are intelligent, borderline human in the way they talk, the way they walk, and the way they interact with their environments. These kinds of robots seem to be creatures that won't exist any time soon, and they probably won't, but with every advancement in the field of robotics and autonomous systems, we step closer to that reality. One key sense that helps a robot communicate and understand its environment is sight, and the area of robotics that strives towards the goal of giving machines the sense of sight is computer vision.

1.1 Area of Research

Human Pose Estimation (HPE) is an area of research within computer vision that aims to teach robots how to make sense of the human form and the motions it is capable of performing. It involves the identification and classification of the joints in the human body, capturing a set of coordinates for each joint, known as a key point, that can describe the pose of a person. HPE has a wide set of uses in many fields: In games, with motion capture technologies reliant on HPE, it allows developers to code program more realistic and fluid character movements. In healthcare, healthcare providers can monitor a patient's movements and detect any abnormalities. Augmented reality, allows the user to interact with the digital content in more natural and intuitive ways with gestures. And finally, the use-case that is the primary focus of this project, is sports training. HPE can be used to analyse a user's performance, identify areas for improvement, and develop personalised training programs based on the physical level of the user. For example, HPE could be used to analyse a runner's form, e.g. How straight is their back? What part of the foot are they landing on? Are they leaning more to one side?..., and provide feedback on how to improve their technique. HPE can be used to collect data about any exercises where the movement of the body is vital to its effectiveness.

1.2 Relevance to the Degree Programme

I am enrolled in the MSc programme Robotics and Autonomous Systems at the University of Lincoln through the AgriFoRwArdS CDT. Throughout the first

two semesters, I studied the principles of robotics, artificial intelligence, machine learning, and computer vision. Principles from all of these subjects are applied within the area of research on human pose estimation. As to my affiliation with the AgriFoRwArdS CDT, their focus is on the production and use of AI, ML, and CV applications to help the farming and agricultural industries, as Lincolnshire is an agricultural region of the UK. Human pose estimation has previously been used in agritech applications (Moysiadis et al. 2022), with its ability to facilitate human-robot interactions in the field for fruit picking, robotic carts will follow the worker through the field to hold the produce and take it away once full. HPE allows these robots to understand gesture commands the worker may give it, and gives the robot an understanding of humans that allows it to find and follow them without driving into them.

1.3 Background of the Topic

When computer vision gained popularity in the late 1960s and early 1970s, HPE research had its start. Scientists first focused on basic problems like as shape analysis, object recognition, and visual understanding. As computer vision developed, HPE became a stand-alone area of study (Lynn 2023). Historically, HPE was frequently described probabilistically to account for likely inference ambiguities. Since deep learning has been more widely used, the focus has switched to end-to-end trainable models because of their ability to extract intricate patterns and postures from data. Traditionally, computer vision systems have assessed an object’s or person’s posture by geometric calculations and feature-based techniques. But, the biggest developments in HPE came with the advent of deep neural networks, convolutional neural networks, and computer vision. The field has advanced considerably in spite of these challenges, and more recent techniques that make use of properly designed neural networks may provide amazing results in challenging scenarios involving a large number of, perhaps veiled, interacting individuals (Liu and Yuan 2018). Now that these detections have the necessary technology and are sufficiently precise, they may be employed for commercial purposes. It also offers a wealth of new application potential and signifies a major change in HPE’s overall direction.

2 Aims and Objectives

2.1 Aims

- Specific: The goal of this project is to create an AI fitness personal trainer application capable of using video input to detect the presence of the user, finding and identifying the keypoints on the user’s body, analysing the poses created by the connections between the keypoints (while making inferences on the position of occluded keypoints), and of giving the user recommendations on how to improve the user’s form to help minimise the risk of injury and maximise the effectiveness of the chosen exercise. If there is time, a potential add-on would be to make this a smartphone application for ease of use.

- Mesurable: The success of this application can be measured by its ability to detect certain movements. First by its ability to detect the human shape then the exercises and the form.
- Achievable: This task is certainly achievable, I have previous experience with image processing using python.
- Relevant: This project is worthwhile for many reasons, a few of them I describe in section 2.3. below.
- Timebound: The deadline for this research project is currently set to the 29th of August 2024. Giving a two month timeframe to complete the project.

2.2 Issues to Explore

There are many issues relating to the subject of human pose estimation. Namely, how to detect keypoints on the human body more efficiently to allow for real-time detection and processing. How to effectively infer the positions of keypoints occluded by other body parts or obstacles. And how to use the data to derive recommendations on how to improve the user's form.

2.3 Motivation

Over the last few decades, the percentage of people who are either overweight or obese has steadily increased. This is a concerning trend as being obese increases the risk of many other health conditions such as type 2 diabetes, coronary heart disease, some types of cancer, and stroke (NHS 2023). As seen in Figure 1, over 60% of the UK population is considered overweight or obese.




Fig. 1. Percentage of people who are either obese or overweight, by year (Source: Health Survey for England, 2018)

To help solve that issue, it is important to encourage people to workout, but going to the gym can be daunting, especially when you don't know how to perform exercises properly. Creating an AI personal trainer will give people the knowledge and confidence to exercise more at home or the gym. Furthermore, exercise has many benefits other than weight loss, it reduces the risk of the issues mentioned above and has been shown to improve mental health (NHS 2021).

3 Literature Survey

3.1 BlazePose: On-device Real-time Body Pose tracking (Bazarevsky et al. 2020)

This 2020 Google Research by Valentin Bazarevsky et al. demonstrates a particular type of convolutional neural network architecture called BlazePose which was created specifically for the real-time mobile human posture estimation task. With the ability to build 33 body keypoints for an individual and operating at over 30 frames per second on a Pixel 2, it is perfect for real-time applications such as fitness tracking and sign language recognition. BlazePose’s primary contributions are a lightweight body pose estimate neural network and a novel body pose tracking method. Regression and heatmaps are both used by the network to determine keypoint coordinates, which probably improves efficiency and accuracy. An important development in the realm of on-device real-time body posture monitoring is this technology. It is a useful instrument for a variety of applications due to its efficiency and adaptability. BlazePose has a bright future ahead of it, and it will be interesting to watch how this technology develops and is applied in many real-world scenarios.



Fig. 2. BlazePose results on yoga and fitness poses

3.2 Realtime Multi-Person 2D Pose Estimation using Part Affinity Fields (Cao et al. 2017)

In this 2017 paper by Zhe Cao et al., an effective technique for predicting posture for several people is Real-Time Multi-Person 2D posture estimate using Part Affinity Fields (PAFs) is presented. This method seeks to identify several individuals’ 2D postures inside a picture. The system gains the ability to link body parts with specific persons in the image by using a nonparametric representation called Part Affinity Fields (PAFs). PAFs, or groups of 2D vector fields, represent the limbs’ orientation and posture inside an image. Regardless of the

number of persons in the image, the system’s architecture for a greedy bottom-up parsing phase that produces exceptional accuracy and real-time speed.



Fig. 3. Top: Multi-person pose estimation. Body parts belonging to the same person are linked. Bottom left: Part Affinity Fields (PAFs) corresponding to the limb connecting right elbow and right wrist. The color encodes orientation. Bottom right: A zoomed in view of the predicted PAFs. At each pixel in the field, a 2D vector encodes the position and orientation of the limbs.

It is therefore especially appropriate for use cases that include real-time data. The main contribution of this strategy is the separation of image population and runtime complexity. Compared to previous methods, which often saw increased computational costs as the number of participants increased, this represents a significant breakthrough. In summary, real-time multi-person 2D posture estimation using Part Affinity Fields presents this innovative and efficient approach for multi-person posture estimation. Because it can do real-time inference regardless of the amount of people in the image, it is a helpful tool in its sector.

3.3 Simple Baselines for Human Pose Estimation and Tracking (Xiao, Wu, and Wei 2018)

Simple Baselines for Human Pose Estimation and Tracking by Bin Xiao et al. is a significant addition to the field of human pose estimation. This study provides baseline methods that are practical and easy to apply for coming up with and evaluating new ideas in the sector. The main contribution of this study is the baseline tracking and posture estimation algorithms established. These methods are surprisingly effective, producing state-of-the-art performance on challenging benchmarks.

This approach’s simplicity is one of its main advantages. This work shows that a straightforward approach may still provide outstanding results, even in the face of the growing complexity of algorithms and systems in this sector. For the purpose of developing and assessing new techniques, this simplifies the process of analysing and comparing algorithms. It can be concluded that Simple Baselines




Fig. 4. Some sample results on PoseTrack Challenge test set

for Human Pose Estimation and Tracking offers a productive and successful method for these tasks. Its ease of use and potency make it a useful instrument in its field, and there is a lot of room for expansion.

3.4 Pose2Seg: Detection Free Human Instance Segmentation (Zhang et al. 2019)

Pose2Seg, created by Song-Hai Zhang et al., is an innovative approach to human instance segmentation. This method provides a unique posture-based instance segmentation framework for humans, which divides instances based on human position, as opposed to proposal region identification. Usually, image instance segmentation begins with object detection and proceeds to segment the object from the detection bounding-box. Conversely, Pose2Seg takes into account the uniqueness of the “human” category, which is precisely specified by the posture skeleton. The human posture skeleton may be used to identify instances with severe occlusion instead of bounding-boxes.



Fig. 5. Heavily occluded people are better separated using human pose than using bounding-box.

This method shows that the pose-based framework can better handle occlusion and obtain higher accuracy on the human instance segmentation test compared to the most sophisticated detection-based approach. Additionally, the publication “Occluded Human (OCHuman)” introduces a new benchmark for occluded persons with full annotations, including instance masks, bounding boxes, and human position. This dataset has 8110 meticulously annotated human occurrences dispersed throughout 4731 images. All things considered, Pose2Seg offers

a fresh and effective method for segmenting human instances. Its detection-free segmentation based on human postural capability makes it a top tool in its field with a lot of future application possibilities.

Summary The science of human posture evaluation has greatly benefited from these four publications. BlazePose makes real-time body location tracking of mobile devices easy and effective. To illustrate the effectiveness of nonparametric representations in the precise estimation of multi-person poses, Realtime Multi-Person 2D Pose Estimation uses component affinity fields. The realistic baseline techniques for posture estimation and tracking provided by Simple Baselines for Human posture Estimation and Tracking emphasise the importance of simplicity in obtaining high performance. Pose2Seg offers a unique pose-based instance segmentation framework and is the last illustration of the potential for detection-free human instance segmentation. By building on the concepts of others, proposing fresh ideas, and refining previously established methods, each study advances the subject of human pose assessment.

4 Research Methods

This research project will attempt to provide a framework for performing physical interventions that enhance the physical well-being of patient populations and to extract skeleton-based features that can quantify the physical status of each user. To avoid ethical concerns, there will be no data collection or participant enrollment in this research project. This research will involve the process of analysing numerical data from video frames. This numerical data will refer to the coordinates of the key points and which keypoints are linked. Therefore, this research is quantitative.

This research will act as a study about the creation of an application which provides a framework for sports professionals to set their desired ranges of motion and form constraints on specific exercises to create personalised workout regimens.

The extracted video markers will be analysed to determine the efficacy and accuracy of the pose estimation techniques used. Each individual exercise type will require its own model trained for that specific movement. Video markers will be extracted from recording I will take of myself and the ones provided by my supervisor to evaluate pose estimation accuracy.

The expected outcome is to have many models capable of detecting the user's poses and extracting useful and relevant data for sports professionals, regular gym goers, and home workout enjoyers to understand and use.

5 Ethical Considerations

After careful consideration, my supervisor and I have decided to avoid the need to submit an ethical application by not collecting any data ourselves or involving

any external human participation. All the data used will be either given to me by my supervisor or produced by myself for testing and validation procedures only.

6 Project Plan and Risk Analysis

6.1 Project Plan

Figure 6 is a gantt chart showing the phases of work I will be attempting to follow over the course of this project. The two greyed out bars are periods where I will be unable to work on this project as I'll be travelling on AgriFoRwArdS CDT duties. There are three major milestones for this project, the submission date for this proposal on the 4th of July, the submission date for this project on the 29th of August, and the in person viva defence of the MSc on the week of the 16th of September.

6.2 Risk Analysis

ID	Description	Impact	Likelihood	mitigation strategy
R1	A program doesn't work or is incompatible with my machine	high	medium	find a alternative program
R2	The landmark detection doesn't function for everyone	low	medium	We won't be able to know if that is the case
R3	Too much data is created for my computer to store	high	low	Reduce the amount of data and make smaller models
R4	The movement detections are slow which cause them to be inaccurate	high	medium	Reduce the amount of keypoints to detect to the strictly necessary ones

Table 1. Table describing the potential risks, how severe they are, how likely they are, and what to do if they occur




Fig. 6. Gantt Chart

References

- Bazarevsky, Valentin et al. (2020). *BlazePose: On-device Real-time Body Pose tracking*. arXiv: 2006.10204 [cs.CV]. URL: <https://arxiv.org/abs/2006.10204>.
- Cao, Zhe et al. (2017). *Realtime Multi-Person 2D Pose Estimation using Part Affinity Fields*. arXiv: 1611.08050 [cs.CV]. URL: <https://arxiv.org/abs/1611.08050>.
- Liu, Mengyuan and Junsong Yuan (2018). “Recognizing human actions as the evolution of pose estimation maps”. In: *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 1159–1168.
- Lynn, Trevor (July 2023). *Pose Estimation Algorithms: History and Evolution*. Roboflow Blog. URL: <https://blog.roboflow.com/pose-estimation-algorithms-history/>.
- Moysiadis, Vasileios et al. (2022). “An Integrated Real-Time Hand Gesture Recognition Framework for Human–Robot Interaction in Agriculture”. In: *Applied Sciences* 12.16. ISSN: 2076-3417. DOI: 10.3390/app12168160. URL: <https://www.mdpi.com/2076-3417/12/16/8160>.
- NHS (2021). *Benefits of Exercise*. URL: <https://www.nhs.uk/live-well/exercise/exercise-health-benefits/>.
- (2023). *Obesity*. URL: <https://www.nhs.uk/conditions/obesity/>.
- Xiao, Bin, Haiping Wu, and Yichen Wei (2018). *Simple Baselines for Human Pose Estimation and Tracking*. arXiv: 1804.06208 [cs.CV]. URL: <https://arxiv.org/abs/1804.06208>.
- Zhang, Song-Hai et al. (2019). *Pose2Seg: Detection Free Human Instance Segmentation*. arXiv: 1803.10683 [cs.CV]. URL: <https://arxiv.org/abs/1803.10683>.

Word Count: 2052